

Systems Programming and Computer Architecture

Janis Hutz
<https://janishutz.com>

January 6, 2026

TITLE PAGE COMING SOON

“If you are using CMake to solve the exercises... First off, sorry that you like CMake“

- Timothy Roscoe, 2025

HS2025, ETHZ
Summary of the Lectures and Lecture Slides

Quotes

“An LLM is a lossy index over human statements”

- Professor Buhmann, Date unknown

“If you are using CMake to solve the exercises... First off, sorry that you like CMake”

“You can't have a refrigerator behave like multiple refrigerators”

“Why is C++ called C++ and not ++C? It's like you don't get any value and then it's incremented, which is true”

- Timothy Roscoe, 2025

Contents

1 The C Programming Language	4
1.1 Basics	4
1.1.1 Control Flow	5
1.1.2 Declarations	6
1.1.3 Operators	7
1.1.4 Arrays	8
1.1.5 Strings	8
1.1.6 Integers in C	9
1.2 The C preprocessor	10
2 x86 Assembly	11
3 Hardware	12

1 The C Programming Language

I can clearly C why you'd want to use C. Already sorry in advance for all the bad C jokes that are going to be part of this section

C is a compiled, low-level programming language, lacking many features modern high-level programming languages offer, like Object Oriented programming, true Functional Programming (like Haskell implements), Garbage Collection, complex abstract datatypes and vectors, just to name a few. (It is possible to replicate these using Preprocessor macros, more on this later).

On the other hand, it offers low-level hardware access, the ability to directly integrate assembly code into the .c files, as well as bit level data manipulation and extensive memory management options, again just to name a few.

This of course leads to C performing excellently and there are many programming languages whose compiler doesn't directly produce machine code or assembly, but instead optimized C code that is then compiled into machine code using a C compiler. This has a number of benefits, most notably that C compilers can produce very efficient assembly, as lots of effort is put into the C compilers by the hardware manufacturers.

There are many great C tutorials out there, a simple one (as for many other languages too) can be found [here](#)

1.1 Basics

C uses a very similar syntax as many other programming languages, like Java, JavaScript and many more... to be precise, it is *them* that use the C syntax, not the other way around. So:

File: 00_intro.c

```

1 // This is a line comment
2 /* this is a block comment */
3 #include "01_func.h" // Relative import
4
5 int i = 0; // This allocates an integer on the stack
6
7 int main( int argc, char *argv[] ) {
8     // This is the function body of a function (here the main function)
9     // which serves as the entrypoint to the program in C and has arguments
10    printf( "Argc: %d\n", argc ); // Number of arguments passed, always >= 1
11        // (first argument is the executable name)
12    for ( int i = 0; i < argc; i++ ) // For loop just like any other sane programming language
13        printf( "Arg %d: %s\n", i, argv[ i ] ); // Outputs the i-th argument from CLI
14
15    get_user_input_int( "Select a number" ); // Function calls as in any other language
16    return 0; // Return a POSIX exit code
17 }
```

In C we are referring to the implementation of a function as a (**function**) **definition** (correspondingly, *variable definition*, if the variable is initialized) and to the definition of the function signature (or variables, without initializing them) as the (**function**) **declaration** (or, correspondingly, *variable declaration*).

C code is usually split into the source files, ending in .c (where the local functions and variables are declared, as well as all function definitions) and the header files, ending in .h, usually sharing the filename of the source file, where the external declarations are defined. By convention, no definition of functions are in the .h files, and neither variables, but there is nothing preventing you from putting them there.

File: 01_func.h

```

1 #include <stdio.h> // Import from system path
2 // (like library imports in other languages)
3
4 int get_user_input_int( char prompt[] );
```

1.1.1 Control Flow

Many of the control-flow structures of C can be found in the below code snippet. A note of caution when using goto: It is almost never a good idea (can lead to unexpected behaviour, is hard to maintain, etc). Where it however is very handy is for error recovery (and cleanup functions) and early termination of multiple loops (jumping out of a loop). So, for example, if you have to run multiple functions to set something up and one of them fails, you can jump to a label and have all cleanup code execute that you have specified there. And because the labels are (as in Assembly) simply skipped over during execution, you can make very nice cleanup code. We can also use continue and break statements similarly to Java, they do not however accept labels. (Reminder: continue skips the loop body and goes to the next iteration)

File: 01_func.c

```

1 #include "01_func.h"
2 #include <stdio.h>
3
4 int get_user_input_int( char prompt[] ) {
5     int input_data;
6     printf( "%s", prompt );           // Always wrap strings like this for printf
7     scanf( "%d", &input_data );      // Get user input from CLI
8     int input_data_copy = input_data; // Value copied
9
10    // If statements just like any other language
11    if ( input_data )
12        printf( "Not 0" );
13    else
14        printf( "Input is zero" );
15
16    // Switch statements just like in any other language
17    switch ( input_data ) {
18        case 5:
19            printf( "You win!" );
20            break; // Doesn't fall through
21        case 6:
22            printf( "You were close" ); // Falls through
23        default:
24            printf( "No win" ); // Case for any not covered input
25    }
26
27    while ( input_data > 1 ) {
28        input_data -= 1;
29        printf( "Hello World\n" );
30    }
31
32    // Inversed while loop (executes at least once)
33    do {
34        input_data -= 1;
35        printf( "Bye World\n" );
36        if ( input_data_copy == 0 )
37            goto this_is_a_label;
38    } while ( input_data_copy > 1 );
39
40 this_is_a_label:
41     printf( "Jumped to label" );
42     return 0;
43 }
```

1.1.2 Declarations

We have already seen a few examples for how C handles declarations. In concept they are similar (and scoping works the same) to most other C-like programming languages, including Java.

File: 02_declarations.c

```

1 int my_int;           // Allocates memory on the stack.
2                         // Variable is global (read / writable by entire program)
3 static int my_local_int; // only available locally (in this file)
4 extern const char *var; // Defined in some other file
5 const int MY_CONST = 10; // constant (immutable), convention: SCREAM_CASE
6
7 enum { ONE, TWO } num; // Enum. ONE will get value 0, TWO has value 1
8
9 enum { 0 = 2, T = 1 } n; // Enum with values specified
10
11 // Structs are like classes, but contain no logic
12 struct MyStruct {
13     int el1;
14     int el2;
15 };
16
17 int fun( int j ) {
18     static int i = 0;           // Persists across calls of fun
19     short my_var = 1;          // Block scoped (deallocated when going out of scope)
20     int my_var_dbl = (int) my_var; // Explicit casting (works between almost all types)
21     return i;
22 }
23
24 int main( int argc, char *argv[] ) {
25     if ( ( my_local_int = fun( 10 ) ) ) {
26         // Every c statement is also an expression, i.e. you can do the above!
27     }
28     struct MyStruct test;        // Allocate memory on stack for struct
29     struct MyStruct *test_p = &test; // Pointer to memory where test resides
30     test.el1 = 1;               // Direct element access
31     test_p->el2 = 2;           // Via pointer
32     return 0;
33 }
```

A peculiarity of C is that the bit-count is not defined by the language, but rather the hardware it is compiled for.

C data type	typical 32-bit	ia32	x86-64
char	1	1	1
short	2	2	2
int	4	4	4
long	4	4	8
long long	8	8	8
float	4	4	4
double	4	8	8
long double	8	10/12	16

Table 1.1: Comparison of byte-sizes for each datatype on different architectures

By default, integers in C are signed, to declare an unsigned integer, use `unsigned int`. Since it is hard and annoying to remember the number of bytes that are in each data type, C99 has introduced the extended integer types, which can

be imported from `stdint.h` and are of form `int<bit count>.t` and `uint<bit count>.t`, where we substitute the `<bit count>` with the number of bits (have to correspond to a valid type of course).

Another notable difference of C compared to other languages is that C doesn't natively have a boolean type, by convention a `short` is used to represent it, where any non-zero value means true and 0 means false. Since boolean types are quite handy, the `!` syntax for negation turns any non-zero value of any integer type into zero and vice-versa. C99 has added support for a `bool` type via `stdbool.h`, which however is still an integer.

Notably, C doesn't have a very rigid type system and lower bit-count types are implicitly cast to higher bit-count data types, i.e. if you add a `short` and an `int`, the `short` is cast to `short` (bits 16-31 are set to 0) and the two are added. Explicit casting between almost all types is also supported. Some will force a change of bit representation, but most won't (notably, when casting to and from float-like types, minus to `void`)

Another important feature is that every C statement is also an expression, see above code block for example.

The `void` type has **no** value and is used for untyped pointers and declaring functions with no return value

It is also possible to define a custom type using `typedef <type it represents> <name of the new type>`

1.1.3 Operators

The list of operators in C is similar to the one of Java, etc. In Table 1.2, you can see an overview of the operators, sorted by precedence in descending order. You may notice that the `&` and `*` operators appear twice. The higher precedence occurrence is the address operator and dereference, respectively, and the lower precedence is `bitwise` and `multiplication`, respectively.

Very low precedence belongs to boolean operators `&&` and `||`, as well as the ternary operator and assignment operators

Operator	Associativity
<code>() [] -> .</code>	Left-to-right
<code>! ~ ++ -- + - * & (type) sizeof</code>	Right-to-left
<code>* / %</code>	Left-to-right
<code>+ -</code>	Left-to-right
<code><< >></code>	Left-to-right
<code>< <= >= ></code>	Left-to-right
<code>== !=</code>	Left-to-right
<code>& (logical and)</code>	Left-to-right
<code>^ (logical xor)</code>	Left-to-right
<code> (logical or)</code>	Left-to-right
<code>&& (boolean and)</code>	Left-to-right
<code> (boolean or)</code>	Left-to-right
<code>? : (ternary)</code>	Right-to-left
<code>= += -= *= /= %= &= ^= == <= >=</code>	Right-to-left
<code>,</code>	Left-to-right

Table 1.2: C operators ordered in descending order by precedence

Associativity

- Left-to-right: $A + B + C \mapsto (A + B) + C$
- Right-to-left: $A += B += C \mapsto (A += B) += C$

As it should be, boolean and, as well as boolean or support early termination.

The ternary operator works as in other programming languages `result = expr ? res_true : res_false;`

As previously touched on, every statement is also an expression, i.e. the following works

```
printf("%s", x = foo(y)); // prints output of foo(y) and x has that value
```

Pre-increment (`++i`, new value returned) and post-increment (`i++`, old value returned) are also supported by C.

C has an `assert` statement, but do not use it for error handling. The basic syntax is `assert(expr);`

1.1.4 Arrays

C compiler does not do any array bound checks! Thus, always check array bounds. Unlike some other programming languages, arrays are **not** dynamic length.

The below snippet includes already some pointer arithmetic tricks. The variable `data` is a pointer to the first element of the array.

File: 03_arrays.c

```

1 #include <stdint.h>
2 #include <stdio.h>
3
4 int main( int argc, char *argv[] ) {
5     int data[ 10 ];           // Initialize array of 10 integers
6     data[ 5 ] = 5;           // element 5 is now 5
7     *data = 10;              // element 0 is now 5
8     printf( "%d\n", data[ 0 ] ); // print element 0 (prints 10)
9     printf( "%d\n", *data );   // equivalent as above
10    printf( "%d\n", data[ 5 ] ); // print element 5 (prints 5)
11    printf( "%d\n", *( data + 5 ) ); // equivalent as above
12    int multidim[ 5 ][ 5 ];    // 2-dimensional array
13                                // We can iterate over it using two for-loops
14    int init_array[ 2 ][ 2 ] = {
15        {1, 2},
16        {3, 4}
17    };                         // We can initialize an array like this
18    int empty_arr[ 4 ] = {}; // Initialized to 0
19    return 0;
20 }
```

1.1.5 Strings

C doesn't have a `string` data type, but rather, strings are represented (when using ASCII) as `char` arrays, with length of the array $n + 1$ (where n is the number of characters of the string). The extra element is the termination character, called the `null character`, denoted `\0`. To determine the actual length of the string (as it may be padded), we can use `strlen(str, maxlen)` from `string.h`.

File: 04_strings.c

```

1 #include <stdio.h>
2 #include <string.h>
3
4 int main( int argc, char *argv[] ) {
5     char hello[ 6 ] = "hello";           // Using double quotes
6     char world[ 6 ] = { 'w', 'o', 'r', 'l', 'd', '\0' }; // As array
7
8     char src[ 12 ], dest[ 12 ];
9     strcpy( src, "ETHZ", 12 );          // Copy strings (extra elements will be set to \0)
10    strcpy( dest, src, 12 );            // Copy strings (last arg is first n chars to copy)
11    if ( strcmp( src, dest, 12 ) ) // Compare two strings. Returns 1 if src > dest
12        printf( "Hello World" );
13    strcat( dest, " is in ZH", 12 ); // Concatenate strings
14    return 0;
15 }
```

1.1.6 Integers in C

As a reminder, integers are encoded as follows in big endian notation, with x_i being the i -th bit and w being the number of bits used to represent the number:

- **Unsigned:** $\sum_{i=0}^{w-1} x_i \cdot 2^i$
- **Signed:** $-x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-1} x_i \cdot 2^i$ (two's complement notation, with x_{w-1} being the sign-bit)

The minimum number representable is 0 and -2^{w-1} , respectively, whereas the maximum number representable is $2^w - 1$ and $2^{w-1} - 1$

We can use the shift operators to multiply and divide by two. Shift operations are usually *much* cheaper than multiplication and division

1.2 The C preprocessor

To have gcc stop compilation after running through cpp, the C preprocessor, use `gcc -E <file name>`.

Imports in C are handled by the preprocessor, that for each `#include <file1.h>`, the preprocessor simply copies the contents of the file recursively into one file.

Depending on if we use `#include <file1.h>` or `#include "file1.h"` the preprocessor will search for the file either in the system headers or in the project directory. Be wary of including files twice, as the preprocessor will recursively include all files (i.e. it will include files from the files we included)

The C preprocessor gives us what are called **preprocessor macros**, which have the format `#define NAME SUBSTITUTION`.

File: 00_macros.c

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #define FOO      BAZ
4 #define BAR( x ) ( x + 3 )
5 #define SKIP_SPACES( p )
6     do {
7         while ( p > 0 ) { p--; }
8     } while ( 0 )
9 #define COMMAND( c ) { #c, c##_command } // Produces { "<val(c)>", "<val(c)>_command" }
10
11 #ifdef FOO // If macro is defined, ifndef for if not defined
12     #define COURSE "SPCA"
13 #else
14     #define COURSE "Systems Programming and Computer Architecture"
15 #endif
16
17 #if 1
18     #define OUT HELLO // if statement
19 #endif
20
21 int main( int argc, char *argv[] ) {
22     int i = 10;
23     SKIP_SPACES( i );
24
25     printf( "%s", COURSE );
26
27     return EXIT_SUCCESS;
28 }
```

To avoid issues with semicolons at the end of preprocessor macros that wrap statements that cannot end in semicolons, we can use a concept called semicolon swallowing. For that, we wrap the statements in a `do ... while(0)` loop, which is removed by the compiler on compile, also taking with it the semicolon.

There are also a number of predefined macros:

- `__FILE__`: Filename of processed file
- `__LINE__`: Line number of this usage of macro
- `__DATE__`: Date of processing
- `__TIME__`: Time of processing
- `__STDC__`: Set if ANSI Standard C compiler is used
- `__STDC_VERSION__`: The version of Standard C being compiled
- ... many more

In headers, we typically use `#ifndef __FILENAME_H__` followed by a `#define __FILENAME_H__` or the like to check if the header was already included before

2 x86 Assembly

3 Hardware