

Systems Programming and Computer Architecture

Robin Bacher, Janis Hutz
<https://github.com/janishutz/eth-summaries>

January 12, 2026

TITLE PAGE COMING SOON

“If you are using CMake to solve the exercises... First off, sorry that you like CMake“

- Timothy Roscoe, 2025

HS2025, ETHZ
Summary of the Lectures and Lecture Slides

Quotes

“An LLM is a lossy index over human statements”

- Professor Buhmann, Date unknown

“If you are using CMake to solve the exercises... First off, sorry that you like CMake”

“You can't have a refrigerator behave like multiple refrigerators”

“Why is C++ called C++ and not ++C? It's like you don't get any value and then it's incremented, which is true”

- Timothy Roscoe, 2025

Contents

1	Introduction	4
2	x86 Assembly	5
2.1	The syntax	6
2.1.1	Registers	6
2.1.2	Instructions	6
2.2	Data types	7
2.2.1	Arrays	7
2.2.2	Structures	7
2.2.3	Nested / Multidimensional arrays, Struct arrays	7
2.2.4	Multi-Level arrays	8
2.2.5	Unions	8
2.3	Operations	9
2.3.1	Arithmetic Operations	9
2.3.2	Condition Codes	9
2.3.3	Jumping	10
2.4	Control Flow	11
2.4.1	Conditional statements	11
2.4.2	While Loops	11
2.4.3	For Loops & Switch	12
2.5	The Stack	13
2.5.1	Calling Conventions	13
3	The C Programming Language	14
3.1	Basics	14
3.1.1	Control Flow	15
3.1.2	Declarations	16
3.1.3	Operators	18
3.1.4	Arrays	19
3.1.5	Strings	19
3.1.6	Integers in C	20
3.1.7	Pointers	21
3.2	The C preprocessor	23
3.3	Memory	24
3.3.1	Dynamic Memory Allocation	25
3.3.2	Garbage Collection	26
3.3.3	Common pitfalls	27
3.4	Variadic functions	28
4	The gcc toolchain	29
5	Hardware	30

1 Introduction

This summary tries to summarize everything that is important to know for this course. It aims to be a full replacement for the slides, but as with all my summaries, there may be missing or incorrect information in here, so use at your own risk. You have been warned!

The summary does *not* follow the order the lecture does. This is to make related information appear more closely to each other than they have in the lecture and the summary assumes you have already seen the concepts in the lectures or elsewhere (or are willing to be thrown in the deep end).

The target semester for this summary is HS2025, so there might have been changes in your year. If there are changes and you'd like to update this summary, please open a pull request in the summary's repo at

<https://github.com/janishutz/eth-summaries>

2 x86 Assembly

Definition: (*Architecture*) Also known as ISA (Instruction Set Architecture) is “The parts of a processor design that one needs to understand to write assembly code”. It includes for example the definition of instructions (and their options) and what registers are available. Notable examples are x86, RISC-V (this one is open-source!), MIPS, ARM, etc

x86-Assembly files usually use the file extension .s or .asm

Definition: (*Microarchitecture*) The implementation of the ISA. It defines the actual hardware layout and how the individual instructions are actually implemented and thus also defines things such as core frequency, cache layout and more.

Thus, the ISA is more or less precisely on the boundary of the software/hardware interface.

Definition: Complex Instruction Set (CISC):

- Stack oriented instruction set: Uses it to pass arguments, save program counter and features explicit push and pop instructions for the stack.
- Arithmetic instructions can access memory
- Condition codes set side effect of arithmetic and logical instructions.
- Design Philosophy: Add new instructions for typical tasks.

Definition: Reduced Instruction Set (RISC):

- Fewer, simpler instructions, commonly with fixed-size encoding. As a result, we might need more to get a given task done. On the other hand, we can execute them with small and fast hardware
- Register-oriented instruction set with many more registers that are used for arguments, return pointers, temporaries, etc.
- Load-Store architecture, i.e. only load and store instructions can access memory
- Thus: No Condition codes

What to choose? Both have advantages that the other has as disadvantage: Compiling for CISC is usually easier and usually results in smaller code size. For RISC however, compiler optimization can give a huge performance uplift and it can run fast with even a simple chip design.

Today, the choices are made based on outside constraints usually. For desktops and servers, there is enough compute to make anything run fast. For embedded systems though, the reduced complexity of RISC makes more sense, but for how long still?

What matters most today are non-technical factors such as existence of code for one ISA or licensing costs (and of course, Geopolitics)

Example A C function that simply adds 2 arguments might be compiled (unoptimized) to this:

File: 01_sum.s

```

1 sum: # Label
2         endbr64          # Indirect Branch target (No effect on code's logic)
3         pushq %rbp        # Preserve caller stack frame
4         movq %rsp, %rbp    # Set up new stack frame
5         movl %edi, -20(%rbp) # Arg1 register -> stack
6         movl %esi, -24(%rbp) # Arg2 register -> stack
7         movl -20(%rbp), %edx # stack[Arg1] -> gp register
8         movl -24(%rbp), %eax # stack[Arg2] -> gp register
9         addl %edx, %eax     # Add Arg1 + Arg2 -> gp register
10        movl %eax, -4(%rbp) # result -> stack
11        movl -4(%rbp), %eax # stack[result] -> return register
12        popq %rbp          # Restore caller stack frame
13        ret                # jump back to caller

```

2.1 The syntax

There are two common styles: AT&T syntax (common on UNIX) and Intel syntax (common on Windows)

The state that is visible to us is:

- PC (Program Counter) that contains the address of the next instruction
- Register file that contains the most used program data
- Condition codes that store status information about most recent arithmetic operation and are used for conditional branching

To view what C code looks like in assembly, we can use `gcc -O0 -S code.c`, which produces `code.s` which contains assembly code.

2.1.1 Registers

x86 assembly is a bit particular with register naming (register names all start in %). The initial 16-bit version of x86 had the following registers (sub registers are registers that can be used to access the high (h suffix) or low (l suffix) half of the register. Only registers ending in x feature these sub registers. They, as well as %si and %di are general purpose):

Name	Sub-registers	Description
%ax	%ah, %al	accumulate
%cx	%ch, %cl	counter
%dx	%dh, %dl	data
%bx	%bh, %bl	base
%si	-	Source index
%di	-	Destination index
%sp	-	Stack pointer
%bp	-	Base pointer
%ip	-	Instruction pointer
%sr	-	Status (flags)

When the architecture was extended to 32-bit, all registers previously available were retained and a 32 bit version of each was introduced with the prefix e. In other words, any 16 bit code would still work as previously, as e.g. the %ax register was simply now the lower 16 bits of the %eax register.

The same happened again when extending to 64-bit, only this time the r prefix was used. So, the register \$eax was now the lower 32 bits of %rax.

Additionally, the following registers are also available, with X to be substituted with 8 through 15: %rX and the lower 32 bits %rXd

2.1.2 Instructions

Instructions usually have a 3 letter mnemonic with a one letter postfix that indicates the number of bytes. The following postfixes are available: b (byte, 1 byte), w (word, 2 bytes), l (long word, 4 bytes) and q (quad, 8 bytes).

The following options can be passed for source and destination: Registers,

Immediates To use a constant value (aka Immediate) in an instruction, we prefix the number with \$ (following number is decimal). To use hex, we can use \$0x, etc.

Memory addresses To treat a register as a memory address, use parenthesis, e.g. (%rax) interprets the value of %rax as a memory address. The instruction will then read the number of bytes, as specified by the postfix of the instruction.

The full syntax for memory address modes is D(Rb, Ri, S), where

- D: Displacement (constant offset), can be 0, 1, 2 or 4 bytes (not bits, if you are confused as I was)
- Rb: Base register (to which offsets, etc are added). Can be any of the 16 integer registers
- Ri: Index register: Any, except for %rsp (and %rbp is also rarely used)
- S: Scale factor (1, 2, 4 or 8, to correct offsets)

The computation that happens is the following: $\text{Mem}[\text{Reg}[Rb] + S * \text{Reg}[Ri] + D]$. Using the `lea src, dest` instruction, we can get the address computed into the dest register. Can be abused for similar arithmetic expressions.

2.2 Data types

Assembly supports the following integer types (where GAS stands for GNU Assembly). If they are signed or unsigned does not matter (as we have seen), so it's up to you to interpret them as one or the other

Intel	GAS	Bytes	C
byte	b	1	[unsigned] char
word	w	2	[unsigned] short
double word	l	4	[unsigned] int
quad word	q	8	[unsigned] long

These integer types are also used for pointer addresses. Assembly also supports floating point numbers. They are stored and operated on in floating point registers.

Intel	GAS	Bytes	C
single	s	4	float
double	d	8	double
extended	t	16	long double

Assembly does not support any aggregate types (such as arrays, structs, etc) natively. You can however (obviously) make your own. In the following section we will cover how C datatypes are compiled into assembly.

2.2.1 Arrays

Arrays of type T and length L are allocated as a contiguous region of memory with size $L * \text{sizeof}(T)$ bytes. We then also store a reference / identifier A to the array (i.e. similar to variable name in C), that holds the address of the first element of the array and can then be used in conjunction with "assembly pointer arithmetic".

Array loops that are written as for-loops in code are usually transformed into do-while loops by the compiler to save one condition check in the beginning, except of course, it might be possible that the loops is never executed.

2.2.2 Structures

We again allocate a contiguous region of memory. Only now, the number of bytes required isn't as straight forward to compute anymore, it is still not hard: We simply sum up the sizes of all members and that will be our required sizes, so for the n members x_i of struct my_struct, we have $\text{sizeof}(\text{my_struct}) = \sum_{i=0}^{n-1} \text{sizeof}(x_i)$.

However, the size of a struct may be different to fulfill alignment requirements set forth by the ISA or operating system. This could mean that the struct takes $n * K$ bytes, where K is the alignment of the largest element

For alignment on x86-64 we have:

- 1 byte (no restrictions)
- 2 bytes (LSB must be 0)
- 4 bytes (2 LSB must be 00)
- 8 bytes (3 LSB must be 000)
- 16 bytes (4 LSB must be 0000)

Another issue is accessing members. The solution to this is however easy and efficient, as at compile time, the offsets are pre-determined and compiled into the setter and/or getter code for the struct.

2.2.3 Nested / Multidimensional arrays, Struct arrays

All of these arrays have similar underlying concepts in the way they are allocated, yet all are a bit different

Common ideas Each of the array's elements are allocated in contiguous regions of memory, with the elements also in contiguous regions of memory. (Imagine it as lining up all elements on a band, i.e. as going through the array in a nested loop and printing all the elements into a single line.) The size of the array is determined by $n * \text{sizeof}(T)$, where T is the type of the elements of the array (or outer array). This is what is different for the lot (as well as accessing elements):

Nested array T is another array. We thus have a recursive definition, where $\text{sizeof}(T)$ resolves to $n * \text{sizeof}(T_1)$, etc. Accessing element i, j, k is handled as follows: $o = i * \text{sizeof}(T) + j * \text{sizeof}(T_1) + k * \text{sizeof}(T_2)$, with T_1 and T_2 the types of the nested arrays

Struct arrays T is a struct.

2.2.4 Multi-Level arrays

In comparison to multidimensional arrays, we have arrays of pointers that contain either more arrays of pointers, (normal) arrays or pointers to other data types. The size of such a Multi-Level array is determined by: $n * \text{sizeof}(\text{ptr})$, where `sizeof(ptr)` is the platform-specific size of a pointer and n is the number of elements in the array

To do an access, we need to do two (or more) memory reads, which we can again do using address computations.

The benefit of these kinds of arrays is that we can store arbitrary data types together in an array, giving us more flexibility.

2.2.5 Unions

Since unions can hold any of the elements listed (but only one at a time), we allocate based on the size of the largest element.

2.3 Operations

Assembly operations include performing arithmetic or logic functions on registers or memory data, transferring data between memory and registers and transferring control (conditional or unconditional jumps).

Note that move instructions *cannot* move data directly from memory to memory.

The following instruction formats are commonly used:

File: 00_operations.s

```

1 main:
2     movq %rax, %rdx # rax is src, rdx is dest
3     cmp  %rax, %rdx # rax is src2, rdx is src1
4     jmp  func      # func is label
5     notq %rax       # rax is dest (and src)
6     ret  # No argument

```

2.3.1 Arithmetic Operations

Arithmetic / logic operations need a size postfix (replace X below with b (1B), w (2B), l (4B) or q (8B)).

Mnemonic	Format	Computation
addX	Src, Dest	Dest \leftarrow Dest + Src
subX	Src, Dest	Dest \leftarrow Dest - Src
imulX	Src, Dest	Dest \leftarrow Dest * Src
salX	Src, Dest	Dest \leftarrow Dest << Src
sarX	Src, Dest	Dest \leftarrow Dest >> Src (arithmetic)
shrX	Src, Dest	Dest \leftarrow Dest >> Src (logical)
xorX	Src, Dest	Dest \leftarrow Dest ^ Src
andX	Src, Dest	Dest \leftarrow Dest & Src
orX	Src, Dest	Dest \leftarrow Dest Src
incX	Dest	Dest \leftarrow Dest + 1
decX	Dest	Dest \leftarrow Dest - 1
negX	Dest	Dest \leftarrow -Dest
notX	Dest	Dest \leftarrow ~Dest

2.3.2 Condition Codes

Any arithmetic operation (that is truly part of the arithmetic operations group, so not including lea for example) implicitly sets the **condition codes**. The following condition codes were covered in the lecture (operation: $t = a + b$):

- CF (Carry Flag): Set if carry out from MSB (unsigned overflow)
- ZF (Zero Flag): Set if $t == 0$
- SF (Sign Flag): Set if $(a - b) < 0$ (for signed)
- OF (Overflow Flag): Set if two's complement overflow (i.e. $(a>0 \&& b>0 \&& t<0) \mid\mid (a<0 \&& b<0 \&& t>=0)$)

Explicit computation In the below explanations, we always assume $src2 = b$ and $src1 = a$

To explicitly compute them, we can use the `cmpX src2, src1` instruction (with X again any of the size postfixes), that essentially computes $(a - b)$ without setting a destination register. When we execute that instruction, CF is set if $a < b$ (unsigned), ZF is set if $a == b$, SF is set if $a < b$ (signed) and OF is set as above, where $t = a - b$.

Another instruction that is used is `testX src2, src1` (X again a size postfix), and acts like computing $a \& b$ and where ZF is set if $a \& b == 0$ and SF is set if $a \& b < 0$.

Zeroing register We can use a move instruction, but that is less efficient than using `xorl reg, reg`, where `reg` is the 32-bit version of the `reg` we want to zero.

Reading condition codes To read condition codes, we can use the `setC` instructions, where the C is to be substituted by an element of table 1

2.3.3 Jumping

To jump, use `jmp <label>` (unconditional jump) or the jC instructions, with C from table 1

setX	Condition	Description
e	ZF	Equal / Zero
ne	\sim ZF	Not Equal / Not Zero
s	SF	Negative
ns	\sim SF	Nonnegative
g	$\sim(SF \wedge OF) \& \sim ZF$	Greater (signed)
ge	$\sim(SF \wedge OF)$	Greater or equal (signed)
l	SF \wedge OF	Less (signed)
le	$(SF \wedge OF) \mid ZF$	Less or equal (signed)
a	$\sim CF \& \sim ZF$	Above (unsigned)
b	CF	Below (unsigned)

Table 1: Condition code postfixes for jump and set instructions

Conditional Moves

Similar to jC, the same postfixes can be applied to cmovC, for example:

```
cmp    %eax, %edx
cmovle %edx, %eax
```

Will move %edx into %eax, only if %edx is less or equal (le) %eax.

This can be used to, for example, compile ternary expressions from C to Assembly. However, this requires evaluating both possible expressions, which may lead to a (significant) performance overhead.

2.4 Control Flow

Control flow structures from C like `if/else` or `for` are compiled into assembly mainly using jumps and conditional `move`.

By the nature of Assembly and thanks to compilers optimizing aggressively, there is no *single* definitive translation of the C control structures: The compiler may translate it very differently depending on the context of the program.

2.4.1 Conditional statements

A function using an `if/else` construct to choose the maximum of 2 numbers might compile like this:

File: 02_max.s

```

1 max:
2     cmpl %esi, %edi      # Set condition flags
3     jle .IF               # Conditional jump if %edi <= %esi
4     movl %edi, %edx      # %edi -> return register
5     jmp .ELSE
6 .IF:
7     movl %esi, %edx      # %esi -> return register
8 .ELSE:
9     ret

```

A function computing the absolute difference $|x - y|$ using an `if/else` construct, might use a conditional `move` instead:

File: 03_absdiff.s

```

1 absdiff:
2     movl %edi, %eax
3     subl %esi, %eax      # arg2 - arg1 -> eax
4     movl %esi, %edx
5     subl %edi, %edx      # arg1 - arg2 -> edx
6     cmpl %esi, %edi      # Set condition flags
7     cmovle %edx, %eax    # edx -> eax, only if eax <= edx
8     ret

```

2.4.2 While Loops

A recursive factorial function using a `do while` loop may be compiled like this:

File: 04_factorial.s

```

1 factorial:
2     movl $1, %eax        # Setup
3 .AGAIN:
4     imull %edi, %eax    # %eax *= %edi
5     subl $1, %edi        # %edi--
6     cmpl $1, %edi        # 1 < %edi ?
7     jg    .AGAIN          # go back, if yes
8     ret

```

The same function, using a `while` loop instead may lead to this:

File: 05_factorial.s

```
1 factorial:  
2     jmp .COMP          # Check condition  
3 .LOOP:  
4     imull %edx, %eax    # %eax *= %edx  
5     decl %edx           # %edx--  
6 .COMP:  
7     cmpl $1, %edx        # 1 < %edx ?  
8     jg .LOOP            # if yes, go to loop  
9     ret
```

2.4.3 For Loops & Switch

`for` loops follow the same idea as `while` loops, albeit with a few more jumps.

`switch` statements are implemented differently depending on size and case values: Sparse switch statements are compiled as decision trees, whereas large switch statements may become *jump tables*.

These jump tables are usually stored in the `.rodata` section with 8-byte alignment. The jump uses offsets:

```
jmp *.LABEL(, %rsi, 8)
```

and we jump to the effective address of `.LABEL + rsi * 8`

2.5 The Stack

In the below two, we can do this with $x = 1, 2, 4, 8$, each corresponding to a size prefix that is set with X

Stack push `pushX src`: Fetch operand at `src`, decrement `%rsp` by x , then writes the operand at address of `%rsp`

Stack pop `popX dest`: Fetch operand at address of `%rsp`, increment `%rsp` by x , then writes the operand into `dest`

Procedure call / return Use `call LABEL`. This pushes return label to the stack and jumps to the `LABEL`. After this instruction, we also may use the `pushX` instruction to store further registers. Just remember to pop in the correct order with the correct size again!

The `ret` instruction is the return instruction and it will jump back to the caller and execution will continue there.

2.5.1 Calling Conventions

The callee is the function that is called and the caller is the code / function that calls the function.

- `%rax` and `%eax` can be used without first saving (usually used as return)
- Argument registers are caller saved (or not if not needed again)
- `%rsp` should not be modified anyway
- `%rbp` is callee saved and is used as frame pointer (usually set to equal `%rsp` at start of procedure and can be used to access elements of the frame (as it should not change during execution of the function and should always point to the start of the frame))

Name	Description
<code>%rax</code>	Return value, #variable args
<code>%rbx</code>	Base pointer, Callee saved
<code>%rcx</code>	Argument 4
<code>%rdx</code>	Argument 3 (and return 2)
<code>%rsi</code>	Argument 2
<code>%rdi</code>	Argument 1
<code>%rsp</code>	Stack pointer
<code>%rbp</code>	Frame pointer, Callee saved

Name	Description
<code>%r8</code>	Argument 5
<code>%r9</code>	Argument 6
<code>%r10</code>	Static chain pointer
<code>%r11</code>	Temporary
<code>%r12</code>	Callee saved
<code>%r13</code>	Callee saved
<code>%r14</code>	Callee saved
<code>%r15</code>	GOT pointer, callee saved

If we have more than 6 arguments to be passed, we can use the stack for this. If we can do all accesses to the stack relative to the stack pointer, we do not need to update `%rbp` and not even `%rbx`, or we can use it for other purposes.

We can also allocate the entire stack frame immediately by incrementing the stack pointer to the final position and then store data relative to it. To deallocate a stack frame, simply increment the stack pointer.

3 The C Programming Language

I can clearly C why you'd want to use C. Already sorry in advance for all the bad C jokes that are going to be part of this section

C is a compiled, low-level programming language, lacking many features modern high-level programming languages offer, like Object Oriented programming, true Functional Programming (like Haskell implements), Garbage Collection, complex abstract datatypes and vectors, just to name a few. (It is possible to replicate these using Preprocessor macros, more on this later).

On the other hand, it offers low-level hardware access, the ability to directly integrate assembly code into the .c files, as well as bit level data manipulation and extensive memory management options, again just to name a few.

This of course leads to C performing excellently and there are many programming languages whose compiler doesn't directly produce machine code or assembly, but instead optimized C code that is then compiled into machine code using a C compiler. This has a number of benefits, most notably that C compilers can produce very efficient assembly, as lots of effort is put into the C compilers by the hardware manufacturers.

There are many great C tutorials out there, a simple one (as for many other languages too) can be found [here](#)

3.1 Basics

C uses a very similar syntax as many other programming languages, like Java, JavaScript and many more... to be precise, it is *them* that use the C syntax, not the other way around. So:

File: 00_intro.c

```

1 // This is a line comment
2 /* this is a block comment */
3 #include "01_func.h" // Relative import
4
5 int i = 0; // This allocates an integer on the stack
6
7 int main( int argc, char *argv[] ) {
8     // This is the function body of a function (here the main function)
9     // which serves as the entrypoint to the program in C and has arguments
10    printf( "Argc: %d\n", argc ); // Number of arguments passed, always >= 1
11                                // (first argument is the executable name)
12    for ( int i = 0; i < argc; i++ ) // For loop just like any other sane programming language
13        printf( "Arg %d: %s\n", i, argv[ i ] ); // Outputs the i-th argument from CLI
14
15    get_user_input_int( "Select a number" ); // Function calls as in any other language
16    return 0; // Return a POSIX exit code
17 }
```

In C we are referring to the implementation of a function as a (**function**) **definition** (correspondingly, *variable definition*, if the variable is initialized) and to the definition of the function signature (or variables, without initializing them) as the (**function**) **declaration** (or, correspondingly, *variable declaration*).

C code is usually split into the source files, ending in .c (where the local functions and variables are declared, as well as all function definitions) and the header files, ending in .h, usually sharing the filename of the source file, where the external declarations are defined. By convention, no definition of functions are in the .h files, and neither variables, but there is nothing preventing you from putting them there.

File: 01_func.h

```

1 #include <stdio.h> // Import from system path
2                                // (like library imports in other languages)
3
4 int get_user_input_int( char prompt[] );
```

3.1.1 Control Flow

Many of the control-flow structures of C can be found in the below code snippet. A note of caution when using goto: It is almost never a good idea (can lead to unexpected behaviour, is hard to maintain, etc). Where it however is very handy is for error recovery (and cleanup functions) and early termination of multiple loops (jumping out of a loop). So, for example, if you have to run multiple functions to set something up and one of them fails, you can jump to a label and have all cleanup code execute that you have specified there. And because the labels are (as in Assembly) simply skipped over during execution, you can make very nice cleanup code. We can also use continue and break statements similarly to Java, they do not however accept labels. (Reminder: continue skips the loop body and goes to the next iteration)

File: 01_func.c

```

1 #include "01_func.h"
2 #include <stdio.h>
3
4 int get_user_input_int( char prompt[] ) {
5     int input_data;
6     printf( "%s", prompt );           // Always wrap strings like this for printf
7     scanf( "%d", &input_data );       // Get user input from CLI
8     int input_data_copy = input_data; // Value copied
9
10    // If statements just like any other language
11    if ( input_data )
12        printf( "Not 0" );
13    else
14        printf( "Input is zero" );
15
16    // Switch statements just like in any other language
17    switch ( input_data ) {
18        case 5:
19            printf( "You win!" );
20            break; // Doesn't fall through
21        case 6:
22            printf( "You were close" ); // Falls through
23        default:
24            printf( "No win" ); // Case for any not covered input
25    }
26
27    while ( input_data > 1 ) {
28        input_data -= 1;
29        printf( "Hello World\n" );
30    }
31
32    // Inversed while loop (executes at least once)
33    do {
34        input_data -= 1;
35        printf( "Bye World\n" );
36        if ( input_data_copy == 0 )
37            goto this_is_a_label;
38    } while ( input_data_copy > 1 );
39
40 this_is_a_label:
41     printf( "Jumped to label" );
42     return 0;
43 }
```

3.1.2 Declarations

We have already seen a few examples for how C handles declarations. In concept they are similar (and scoping works the same) to most other C-like programming languages, including Java.

File: 02_declarations.c

```

1 int my_int;           // Allocates memory on the stack.
2                         // Variable is global (read / writable by entire program)
3 static int my_local_int; // only available locally (in this file)
4 extern const char *var; // Defined in some other file
5 const int MY_CONST = 10; // constant (immutable), convention: SCREAM_CASE
6
7 enum { ONE, TWO } num; // Enum. ONE will get value 0, TWO has value 1
8
9 enum { O = 2, T = 1 } n; // Enum with values specified
10
11 // Structs are like classes, but contain no logic
12 struct MyStruct {
13     int el1;
14     int el2;
15 };
16
17 // Like structs, but can only hold one of the values!
18 union MyUnion {
19     int ival;
20     float fval;
21     char *sval;
22 };
23
24 int fun( int j ) {
25     static int i = 0;           // Persists across calls of fun
26     short my_var = 1;          // Block scoped (deallocated when going out of scope)
27     int my_var_dbl = (int) my_var; // Explicit casting (works between almost all types)
28     return i;
29 }
30
31 int main( int argc, char *argv[] ) {
32     if ( ( my_local_int = fun( 10 ) ) ) {
33         // Every c statement is also an expression, i.e. you can do the above!
34     }
35     struct MyStruct test;        // Allocate memory on stack for struct
36     struct MyStruct *test_p = &test; // Pointer to memory where test resides
37     struct MyStruct test2;
38     union MyUnion my_uval; // Work exactly like structs for access
39     test.el1 = 1;             // Direct element access
40     test_p->el2 = 2;          // Via pointer
41     test2 = test;             // Copies the struct
42     return 0;
43 }
```

A peculiarity of C is that the bit-count is not defined by the language, but rather the hardware it is compiled for.

C data type	typical 32-bit	ia32	x86-64
char	1	1	1
short	2	2	2
int	4	4	4
long	4	4	8
long long	8	8	8
float	4	4	4
double	4	8	8
long double	8	10/12	16

Table 1: Comparison of byte-sizes for each datatype on different architectures

Type format Be however aware that this table uses the LP64 format for the x86-64 sizes and this is the format all UNIX-Systems use (i.e. Linux, BSD, Darwin (the Mac Kernel)). 64 bit Windows however uses LLP64, i.e. int and long have the same size (32) and long long and pointers are 64 bit.

Integers By default, integers in C are signed, to declare an unsigned integer, use `unsigned int`. Since it is hard and annoying to remember the number of bytes that are in each data type, C99 has introduced the extended integer types, which can be imported from `stdint.h` and are of form `int<bit count>_t` and `uint<bit count>_t`, where we substitute the `<bit count>` with the number of bits (have to correspond to a valid type of course).

Booleans Another notable difference of C compared to other languages is that C doesn't natively have a boolean type, by convention a `short` is used to represent it, where any non-zero value means true and 0 means false. Since boolean types are quite handy, the `!` syntax for negation turns any non-zero value of any integer type into zero and vice-versa. C99 has added support for a `bool` type via `stdbool.h`, which however is still an integer.

Implicit casts Notably, C doesn't have a very rigid type system and lower bit-count types are implicitly cast to higher bit-count data types, i.e. if you add a `short` and an `int`, the `short` is cast to `short` (bits 16-31 are set to 0) and the two are added. Explicit casting between almost all types is also supported. Some will force a change of bit representation, but most won't (notably, when casting to and from float-like types, minus to `void`)

Expressions Every C statement is also an expression, see above code block for example.

Void The `void` type has *no* value and is used for untyped pointers and declaring functions with no return value

Structs Are like classes in OOP, but they contain no logic. We can assign copy a struct by assignment and they behave just like everything else in C when used as an argument for functions in that they are passed by value and not by reference. You can of course pass it also by reference (like any other data type) by setting the argument to type `struct mystruct * name` and then calling the function using `func(&test)` assuming `test` is the name of your struct

Typedef To define a custom type using `typedef <type it represents> <name of the new type>`.

You may also use `typedef` on structs using `typedef struct <struct tag> <name of the new alias>`, you can thus instead of e.g. `struct list_el my_list; write list my_list;`, if you have used `typedef struct list_el list;` before. It is even possible to do this:

```

1 typedef struct list_el {
2     unsigned long val;
3     struct list_el *next;
4 } list_el;
5
6 struct list_el my_list;
7 list_el my_other_list;

```

Namespaces C has a few different namespaces, i.e. you can have the one of the same name in each namespace (i.e. you can have `struct a`, `int a`, etc). The following namespaces were covered:

- Label names (used for `goto`)
- Tags (for `struct`, `union` and `enum`)
- Member names one namespace for each `struct`, `union` and `enum`
- Everything else mostly (types, variable names, etc, including `typedef`)

3.1.3 Operators

The list of operators in C is similar to the one of Java, etc. In Table 2, you can see an overview of the operators, sorted by precedence in descending order. You may notice that the `&` and `*` operators appear twice. The higher precedence occurrence is the address operator and dereference, respectively, and the lower precedence is bitwise and and multiplication, respectively.

Very low precedence belongs to boolean operators `&&` and `||`, as well as the ternary operator and assignment operators

Operator	Associativity
<code>() [] -> .</code>	Left-to-right
<code>! ~ ++ -- + - * & (type) sizeof</code>	Right-to-left
<code>* / %</code>	Left-to-right
<code>+ -</code>	Left-to-right
<code><< >></code>	Left-to-right
<code>< <= >= ></code>	Left-to-right
<code>== !=</code>	Left-to-right
<code>& (logical and)</code>	Left-to-right
<code>^ (logical xor)</code>	Left-to-right
<code> (logical or)</code>	Left-to-right
<code>&& (boolean and)</code>	Left-to-right
<code> (boolean or)</code>	Left-to-right
<code>? : (ternary)</code>	Right-to-left
<code>= += -= *= /= %= &= ^= == <=>=</code>	Right-to-left
<code>,</code>	Left-to-right

Table 2: C operators ordered in descending order by precedence

Associativity

- Left-to-right: $A + B + C \mapsto (A + B) + C$
- Right-to-left: $A += B += C \mapsto (A += B) += C$

As it should be, boolean and, as well as boolean or support early termination.

The ternary operator works as in other programming languages `result = expr ? res_true : res_false;`

As previously touched on, every statement is also an expression, i.e. the following works

```
printf("%s", x = foo(y)); // prints output of foo(y) and x has that value
```

Pre-increment (`++i`, new value returned) and post-increment (`i++`, old value returned) are also supported by C.

C has an `assert` statement, but do not use it for error handling. The basic syntax is `assert(expr);`

3.1.4 Arrays

C compiler does not do any array bound checks! Thus, always check array bounds. Unlike some other programming languages, arrays are **not** dynamic length.

The below snippet includes already some pointer arithmetic tricks. The variable `data` is a pointer to the first element of the array.

File: 03_arrays.c

```

1 #include <stdint.h>
2 #include <stdio.h>
3
4 int main( int argc, char *argv[] ) {
5     int data[ 10 ];           // Initialize array of 10 integers
6     data[ 5 ] = 5;           // element 5 is now 5
7     *data = 10;              // element 0 is now 5
8     printf( "%d\n", data[ 0 ] ); // print element 0 (prints 10)
9     printf( "%d\n", *data );   // equivalent as above
10    printf( "%d\n", data[ 5 ] ); // print element 5 (prints 5)
11    printf( "%d\n", *( data + 5 ) ); // equivalent as above
12    int multidim[ 5 ][ 5 ];    // 2-dimensional array
13                                // We can iterate over it using two for-loops
14    int init_array[ 2 ][ 2 ] = {
15        {1, 2},
16        {3, 4}
17    };                         // We can initialize an array like this
18    int empty_arr[ 4 ] = {}; // Initialized to 0
19    return 0;
20 }
```

3.1.5 Strings

C doesn't have a string data type, but rather, strings are represented (when using ASCII) as char arrays, with length of the array $n + 1$ (where n is the number of characters of the string). The extra element is the termination character, called the null character, denoted `\0`. To determine the actual length of the string (as it may be padded), we can use `strlen(str, maxlen)` from `string.h`.

File: 04_strings.c

```

1 #include <stdio.h>
2 #include <string.h>
3
4 int main( int argc, char *argv[] ) {
5     char hello[ 6 ] = "hello";           // Using double quotes
6     char world[ 6 ] = { 'w', 'o', 'r', 'l', 'd', '\0' }; // As array
7
8     char src[ 12 ], dest[ 12 ];
9     strcpy( src, "ETHZ", 12 );          // Copy strings (extra elements will be set to \0)
10    strcpy( dest, src, 12 );            // Copy strings (last arg is first n chars to copy)
11    if ( strcmp( src, dest, 12 ) ) // Compare two strings. Returns 1 if src > dest
12        printf( "Hello World" );
13    strcat( dest, " is in ZH", 12 ); // Concatenate strings
14    return 0;
15 }
```

3.1.6 Integers in C

As a reminder, integers are encoded as follows in big endian notation, with x_i being the i -th bit and w being the number of bits used to represent the number:

- **Unsigned:** $\sum_{i=0}^{w-1} x_i \cdot 2^i$
- **Signed:** $-x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-1} x_i \cdot 2^i$ (two's complement notation, with x_{w-1} being the sign-bit)

The minimum number representable is 0 and -2^{w-1} , respectively, whereas the maximum number representable is $2^w - 1$ and $2^{w-1} - 1$. `limits.h` defines constants for the minimum and maximum values of different types, e.g. `ULONG_MAX` or `LONG_MAX` and `LONG_MIN`

We can use the shift operators to multiply and divide by two. Shift operations are usually *much* cheaper than multiplication and division. Left shift ($u << k$ in C) always fills with zeros and throws away the extra bits on the left (equivalent to multiplication by 2^k), whereas right shift ($u >> k$ in C) is implementation-defined, either arithmetic (fill with most significant bit, division by 2^k). This however rounds incorrectly, see below) or logical shift (fill with zeros, unsigned division by 2^k).

Signed division using arithmetic right shifts has the issue of incorrect rounding when number is < 0 . Instead, we represent $s/2^k = s + (2^k - 1) >> k$ for $s < 0$ and $s/2^k = s >> k$ for $s > 0$

In expressions, signed values are implicitly cast to unsigned

This can lead to all sorts of nasty exploits (e.g. provide -1 as the argument to `memcpy` and watch it burn, this was an actual exploit in FreeBSD)

Addition & Subtraction

A nice property of the two's complement notation is that addition and subtraction works exactly the same as in normal notation, due to over- and underflow. This also obviously means that it implements modular arithmetic, i.e.

$$\text{Add}_w(u, v) = u + v \bmod 2^w \quad \text{and} \quad \text{Sub}_w(u, v) = u - v \bmod 2^w$$

Multiplication & Division

Unsigned multiplication with addition forms a commutative ring. Again, it is doing modular arithmetic and

$$\text{UMult}_w(u, v) = u \cdot v \bmod 2^w$$

3.1.7 Pointers

On loading of a program, the OS creates the virtual address space for the process, inspects the executable and loads the data to the right places in the address space, before other preparations like final linking and relocation are done.

Stack-based languages (supporting recursion) allocate stack in frames that contain local variables, return information and temporary space. When a procedure is entered, a stack frame is allocated and executes any necessary setup code (like moving the stack pointer, see later). When a procedure returns, the stack frame is deallocated and any necessary cleanup code is executed, before execution of the previous frame continues.

In C a pointer is a variable whose value is the memory address of another variable

Of note is that if you simply declare a pointer using type `* p`; you will get different memory addresses every time. The (Linux)-Kernel randomizes the address space to prevent some common exploits.

File: 05_pointers.c

```

1 #include "01_func.h" // See a few pages up for declarations
2 #include <assert.h>
3 #include <stdio.h>
4 #include <stdlib.h>
5
6 void a_function( int ( *func )( char * ), char prompt[] ) {
7     ( *func )( prompt ); // Call function with arguments
8 }
9
10 int main( int argc, char *argv[] ) {
11     int x = 0;
12     int *p = &x;           // Get x's memory address
13     printf( "%p\n", p ); // Print the address of x
14     printf( "%d\n", *p ); // Dereference pointer (get contents of memory location)
15     *p = 10;              // Dereference assign
16     int **dbl_p = &p;    // Double pointer (pointer to pointer to value)
17     int *null_p = NULL;  // Create NULL pointer
18     *null_p = 1;          // Segmentation fault due to null pointer dereference
19
20     // pointer arithmetic
21     int arr[ 3 ] = { 2, 3, 4 };
22     char c_arr[ 3 ] = { 'A', 'B', 'C' };
23     int *arr_p = &arr[ 1 ];
24     char *c_arr_p = &c_arr[ 1 ];
25     c_arr_p += 1; // Now points to c_arr[2]
26     arr_p -= 1;   // Now points to arr[0]
27
28     char *arr_p_c = (char *) arr_p; // Cast to char pointer (points to first byte of arr[0])
29     printf( "%d", *( arr_p - 5 ) ); // No boundary checks (can access any memory)
30     assert( arr == &( arr[ 0 ] ) ); // Evaluates to true
31     int new_arr[ 3 ] = arr;        // Compile time error (cannot use other array as
32     ↪   initializer)
33     int *new_arr_p = &arr[ 0 ];    // This works
34
35     a_function( &get_user_input_int, c_arr );
36
37     return EXIT_SUCCESS;
}

```

Some pointer arithmetic has already appeared in section 3.1.4, but same kind of content with better explanation can be found here

Pointer Arithmetic Note that when doing pointer arithmetic, adding 1 will move the pointer by `sizeof(type)` bits.

You may use pointer arithmetic on whatever pointer you'd like (as long as it's not a null pointer). This means, you *can* make an array wherever in memory you'd like. The issue is just that you are likely to overwrite something, and that something might be something critical (like a stack pointer), thus you will get **undefined** behaviour! (This is by the way a common concept in C, if something isn't easy to make more flexible (example for `malloc`, if you pass a pointer to memory that is not the start of the `malloc`'d section, you get undefined behaviour), in the docs mention that one gets undefined behaviour if you do not do as it says so... RTFM!)

As already seen in the section arrays (section 3.1.4), we can use pointer arithmetic for accessing array elements. The array name is treated as a pointer to the first element of the array, except when:

- it is operand of `sizeof` (return value is $n \cdot \text{sizeof}(\text{type})$ with n the number of elements)
- its address is taken (then `&a == a`)
- it is a string literal initializer. If we modify a pointer `char *b = "String";` to string literal in code, the "String" is stored in the code segment and if we modify the pointer, we get undefined behaviour

Fun fact : `A[i]` is always rewritten `*(A + i)` by compiler.

Function arguments Another important aspect is passing by value or by reference. You can pass every data type by reference, you can not however pass an array by value (as an array is treated as a pointer, see above).

Body-less loops

```
1 int x = 0;
2 while ( x++ < 10 ); // This is (of course) not a useful snippet, but shows the concept
```

Function pointers A function can be passed as an argument to another function using the typical address syntax with the `&` symbol is annotated as argument using type `(* name)(type arg1, ...)` and is called using `(*func)(arg1, ...)`.

3.2 The C preprocessor

To have gcc stop compilation after running through cpp, the C preprocessor, use `gcc -E <file name>`.

Imports in C are handled by the preprocessor, that for each `#include <file1.h>`, the preprocessor simply copies the contents of the file recursively into one file.

Depending on if we use `#include <file1.h>` or `#include "file1.h"` the preprocessor will search for the file either in the system headers or in the project directory. Be wary of including files twice, as the preprocessor will recursively include all files (i.e. it will include files from the files we included)

The C preprocessor gives us what are called **preprocessor macros**, which have the format `#define NAME SUBSTITUTION`.

File: 00_macros.c

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #define FOO      BAZ
4 #define BAR( x ) ( x + 3 )
5 #define SKIP_SPACES( p )
6     do {
7         while ( p > 0 ) { p--; }
8     } while ( 0 )
9 #define COMMAND( c ) { #c, c##_command } // Produces { "<val(c)>", "<val(c)>_command" }
10
11 #ifdef FOO // If macro is defined, ifndef for if not defined
12     #define COURSE "SPCA"
13 #else
14     #define COURSE "Systems Programming and Computer Architecture"
15 #endif
16
17 #if 1
18     #define OUT HELLO // if statement
19 #endif
20
21 int main( int argc, char *argv[] ) {
22     int i = 10;
23     SKIP_SPACES( i );
24
25     printf( "%s", COURSE );
26
27     return EXIT_SUCCESS;
28 }
```

To avoid issues with semicolons at the end of preprocessor macros that wrap statements that cannot end in semicolons, we can use a concept called semicolon swallowing. For that, we wrap the statements in a `do ... while(0)` loop, which is removed by the compiler on compile, also taking with it the semicolon.

There are also a number of predefined macros:

- `__FILE__`: Filename of processed file
- `__LINE__`: Line number of this usage of macro
- `__DATE__`: Date of processing
- `__TIME__`: Time of processing
- `__STDC__`: Set if ANSI Standard C compiler is used
- `__STDC_VERSION__`: The version of Standard C being compiled
- ... many more

In headers, we typically use `#ifndef __FILENAME_H__` followed by a `#define __FILENAME_H__` or the like to check if the header was already included before

3.3 Memory

In comparison to most other languages, C does not feature automatic memory management, but instead gives us full, manual control over memory. This of course has both advantages and disadvantages.

File: 00_memory.c

```

1 #include <stdlib.h>
2
3 int main( int argc, char *argv[] ) {
4     long *arr = (long *) malloc( 10 * sizeof( long ) ); // Allocate on heap
5     if ( arr == NULL )                                // Check if successful
6         return EXIT_FAILURE;
7     arr[ 0 ] = 5;
8
9     long *arr2;
10    if ( ( arr2 = (long *) calloc( 10, sizeof( long ) ) ) == NULL )
11        return EXIT_FAILURE; // Same as above, but fewer lines and memory zeroed
12
13    // Reallocate memory (to change size). Always use new pointer and do check!
14    if ( ( arr2 = (long *) realloc( arr2, 15 * sizeof( long ) ) ) == NULL )
15        return EXIT_FAILURE;
16
17    free( arr ); // Deallocate the memory
18    arr = NULL; // Best practice: NULL pointer
19    free( arr2 ); // *Can* omit NULLing pointer because end
20
21    return EXIT_SUCCESS;
22 }
```

Notably, the argument `size_t sz` for `malloc`, `calloc` and `realloc` is an `unsigned integer` of some size and differs depending on hardware and software platforms.

`malloc` keeps track of which blocks are allocated. If you give `free` a pointer that isn't the start of the memory region previously `malloc`'d, you get undefined behaviour.

Memory corruption There are many ways to corrupt memory in C. The below code shows off a few of them:

File: 01_mem-corruption.c

```

1 #include <stdlib.h>
2
3 int main( int argc, char **argv ) {
4     int a[ 2 ];
5     int *b = malloc( 2 * sizeof( int ) ), *c;
6     a[ 2 ] = 5;           // assign past the end of an array
7     b[ 0 ] += 2;          // assume malloc zeroes out memory
8     c = b + 3;            // mess up your pointer arithmetic
9     free( &( a[ 0 ] ) ); // pass pointer to free() that wasn't malloc'ed
10    free( b );
11    free( b ); // double-free the same block
12    b[ 0 ] = 5; // use a free()'d pointer
13    // any many more!
14    return 0;
15 }
```

Memory leaks If we allocate memory, but never free it, we use more and more memory (old memory is inaccessible)

Dynamic data structures We build it using structs that have a pointer to another struct inside them. We have to allocate memory for each element and then add the pointer to another struct. For a generic dynamic data structure, make the element a `void pointer`. This in general is the concept used for functions operating on any data type.

3.3.1 Dynamic Memory Allocation

Memory allocated with `malloc` is typically 8- or 16-byte aligned.

Explicit vs. Implicit In explicit memory management, the application does both the allocation *and* deallocation memory, whereas in implicit memory management, the application allocates the memory, but usually a *Garbage Collector* (GC) frees it.

For some languages, like Rust, one would assume that it does implicit allocation, but Rust is a language using explicit management, it's just that the *compiler* and not the programmer decides when to allocate and when to deallocate.

Assumption in this course: Memory is **word** addressed (= 8 Bytes on 64-bit platform).

Goals The allocation should have the highest possible throughput and at the same time the best (i.e. lowest) possible memory utilization. This however is usually conflicting, so we have to balance the two.

Definition: Aggregate payload P_k : All `malloc`'d stuff minus all `free`'d stuff

Definition: Current heap size H_k : Monotonically non-decreasing. Grows when `sbrk` system call is issued.

Definition: Peak memory utilization $U_k = (\max_{i < k} P_i)/H_k$

A bit problem for the `free` function is to know how much memory to free without knowing the size of the to be freed block. This is just one of many other implementation issues:

- (1) How much memory to free? → Headers
- (2) How do we keep track of the free blocks? I.e. where and how large are they? → Free lists
- (3) What do we do with the extra space of a block when allocating a smaller block? → Coalescing
- (4) How do we pick a block? → Placement policies
- (5) How do we reinsert a freed block into the heap? → When to coalesce

This all leads to an issue known as **fragmentation**

Definition: Internal Fragmentation: If for a given block the payload (i.e. the requested size) is smaller than the block size. This depends on the pattern of previous requests and is thus easy to measure

Definition: External Fragmentation: There is enough aggregate heap memory, but there isn't a single large enough free block available. This depends on the pattern of future requests and is thus hard to measure.

Header : Stores size of block and is usually placed in the word that precedes the allocated block (standard method).

Free lists

M1 **Implicit list** using length: Links all blocks and uses a low-order bit to indicate free / allocated, as for aligned blocks, a / some low-order bit(s) are always 0.

M2 **Explicit list** among free blocks using a pointer in the first (and possibly second) word of the block

M3 **Segregated free list**: different free lists for different size classes

M4 **Blocks sorted by size**: Using a balanced tree with pointers within each free block and the length is the key

Definition: Coalescing Connecting two (or more) (free) blocks to form a larger (free) block.

We can do this efficiently in one direction with just a header, however in both directions requires what are referred to as **boundary tags**. They are simply headers on both sides of the block and this allows us to traverse backwards.

We can do coalescing in constant time, by looking at the previous block's footer and next block's header to check if they are free or not.

- If the previous block is free, we can coalesce it by updating its header to include the length of the to be freed block and the middle two words. Same update has to happen to the to be freed block's footer.
- If the next block is free, update its footer to the size of the to be freed block plus the free block's size plus the two words in the middle. Do the same update to the to be freed block's header.

If both blocks are free, then of course we can do this step in one go for both.

Using the headers or boundary tags is just one option to do it and it can be optimized.

Implicit Free List If we use the size that is stored in the header, we know where the next block is going to be already. To know if a block is allocated, we use a low-order bit. This is possible, since if the blocks are aligned, then some of the low-order bits are always 0.

Allocating Blocks: We traverse the list, and choose depending on the placement policy: first fit, next fit or best fit policies. When a block is picked, we might want to split it by adding a header to the remaining part.

Freeing Blocks: We can simply clear the allocated flag, which might lead to fragmentation. Thus: Coalesce the freed blocks.

This is the simplest approach, the price we pay is that allocating is linear w.r.t *all blocks* managed by the allocator.

Explicit Free List To improve on this, here we maintain pointers to the next and preferably previous free block(s). This means that all free blocks form a (doubly) linked list, and we don't traverse already allocated blocks. Since the free blocks aren't used by the program, the list pointers can be stored inside the free blocks.

Allocating Blocks: We remove the element from the list by updating the pointers and we can again use the first fit, next fit or best fit policies.

Freeing Blocks: Requires updating the list pointers, which depends on the list ordering. If boundary tags are used for coalescing, we also have to clear the allocated bit in the footer and header.

The list ordering may be LIFO (last-in-first-out) or address-ordered. LIFO (with first fit) allows constant time freeing, while address-ordered provides better memory utilization (linear time freeing).

To prevent fragmentation, we want to use coalescing again, which can be implemented using boundary tags (again leads to linear time coalescing, if the list is FIFO). This time, the pointers need to be updated too. How this works (and how fast) depends on the list ordering.

Segregated Free List An issue with the previous approaches is finding a *large enough* block, which is why finding a fit remains in linear time. Here, we keep separate lists for different *size classes* to make this faster. For each *size class* a separate free list is maintained. For example, the classes might be:

$$\{1\}, \{2\}, \{3, 4\}, \{5 - 8\}, \dots, \{1025 - 2048\}, \{2049 - 4096\}, \{4097 - \infty\} \quad (\text{Partitions by Powers of 2})$$

Allocating Blocks: We first check the list for the requested size class, and if no fitting block is found, move up to the next list. If we have reached the last list and there is no suitable block, request new memory using `sbrk()`.

Freeing Blocks: Requires adding the newly freed block into the respective free list, potentially after coalescing.

This leads to an increased throughput and better memory utilization as compared to the previous two. Additionally, this structure opens up many optimization options.

The way splitting, coalescing, etc. is implemented varies a lot by implementation.

For example, we might choose to only coalesce if the length of a certain list has reached a certain threshold.

Similarly, if a list of some size class is empty, we might either request new memory and create a new list, or split a block from a higher size class.

Some options used are: *Simple Segregated Storage*, *Segregated Fits*, *Buddy System*

Fun fact : the GNU `malloc` package uses a Segregated Free List, with *Segregated fits*.

Placement policies

- **First fit** Search the list from the beginning, pick first free block that fits. This will usually cause "splinters" at the beginning of the list and can take linear time in the total number of blocks (allocated and free)
- **Next fit** Like first fit, but start at point the previous search finished at. This should be faster, however leads to worse fragmentation.
- **Best fit** Searches the list and chooses the *best* free block that fits and has fewest bytes left over. Leads to lower fragmentation, but is slower than first fit.

3.3.2 Garbage Collection

The memory manager must somehow be able to tell what memory can be freed. In general, we cannot know if memory is going to be used or not, except if there exists no pointer to it anymore. Garbage collectors use graphs to track pointer availability. In other words, a block is reachable if there exists a path from a root node to it.

An easy GC algorithm is called **Mark and Sweep**. It has an extra bit in the header called the *mark bit* and can be built on top of malloc/free. The concept is to use malloc until we “run out of space” and to then run these steps:

- **Mark**: Starts at each root node and sets a mark bit on each reachable block.
- **Sweep**: Scan all blocks and free all blocks that are unmarked.

3.3.3 Common pitfalls

- Dereferencing bad pointers (e.g. passing an `int` to a function expecting a pointer)
- Reading uninitialized memory (memory allocated with `malloc` should be considered garbage)
- Overwriting memory (if you mess up pointer arithmetic or don't do boundary checks)
- Referencing nonexistent variables (variables go out of scope on function returns, except `static`)
- Freeing blocks multiple times (can corrupt the heap)
- Referencing freed blocks (always `NULL` pointers after using `free()`)
- Failing to free blocks (memory leak incoming and make sure to free the ENTIRE data structure!)

Some of these bugs (especially bad references) can usually be found using a debugger.

Substitute `malloc` with a `malloc` that has extra checking code (like UToronto CSRI `malloc` to detect memory leaks)

Another option is using `valgrind` (a memory debugger). Or, simply don't bother with C and use Rust.

3.4 Variadic functions

Variadic functions take a variable number of arguments and use the ... syntax in C. A notable example of such a function is printf

File: 00_variadic.c

```
1 #include <stdarg.h> // Variadic function utilities
2 #include <stdio.h>
3
4 void print_int( unsigned int num_ints, char *msg, ... ) {
5     va_list ap; // keeps track of current argument, similar (in concept) to iterator
6     va_start( ap, msg ); // Initialize the iterator based on last fixed arg
7     for ( int i = 0; i < num_ints; i++ )
8         printf( "int %d = %d\n", i, va_arg( ap, int ) ); // Returns next arg cast to int
9     va_end( ap ); // Free up iterator
10 }
```

4 The gcc toolchain

5 Hardware

Remember: Rust and the like have an `unsafe` block... C's equivalent to this is

```
1 int main( int argc, char *argv[] ) {
2     // Unsafe code goes here
3 }
```

i.e. ***YOU are the one that makes C code safe!***